



Convolutional and Recurrent Neural Networks

Data Science Decal

Hosted by Machine Learning at Berkeley

Agenda

Background

Convolutional Neural Networks

Recurrent Neural Networks

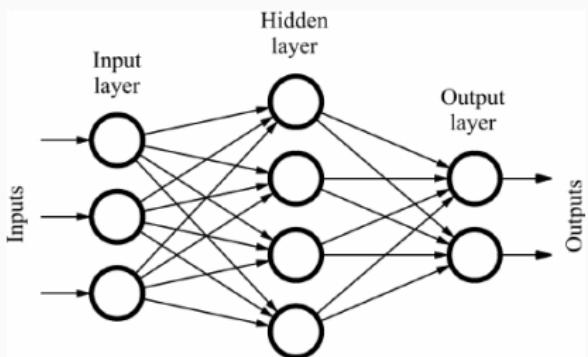
Questions

Demos

Background

Going beyond feed-forward neural networks

- Where we left off:



- Why would such a network not be good for certain tasks?
- How does a neural network deal with complex or high-dimensional data such as images and text?
- ... if we have two hidden layers with 10,000 neurons each,

How are neural networks being used in industry?



Problem: high-dimensional image data



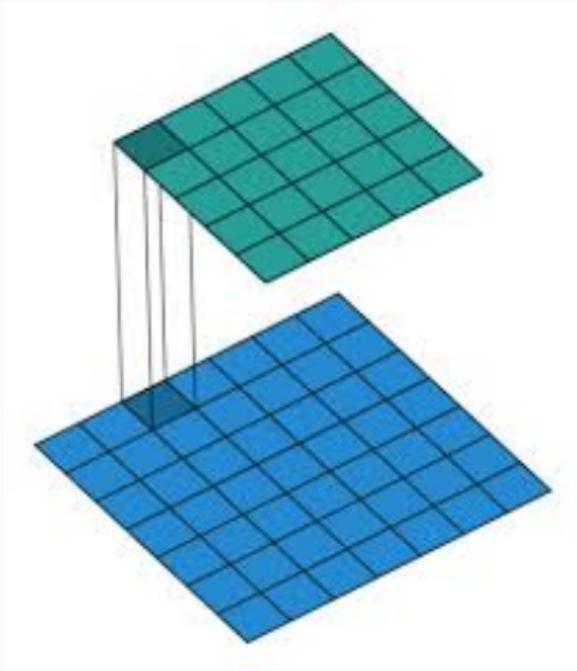
- If we want a neural network to take as input 256×256 images, the input layer will have 65,536 neurons
- ... if we have one hidden layers with 1,000 neurons, we're already at 65 million parameters
- ... if we have two hidden layers with 15,000 neurons each, we're at over a **billion** parameters

Solution: convolutions



- Convolutional architectures allow us to **vastly** reduce the amount of parameters in a network
- Convolutions emulate the response of an individual neuron to visual stimuli
- Convolution operators are translation invariant

Convolution gif



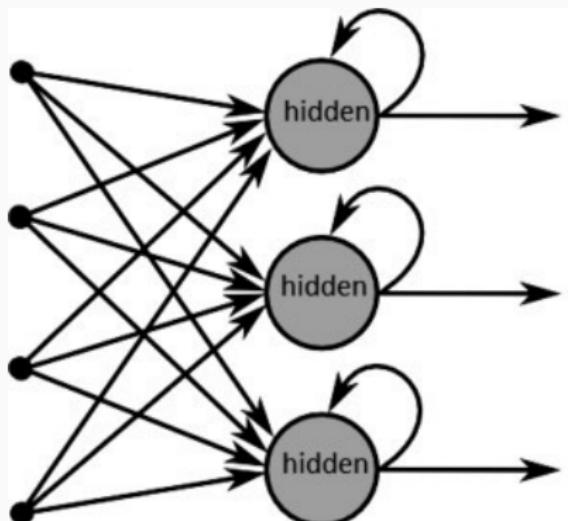
¹<http://iamaaditya.github.io/2016/03/one-by-one-convolution/>

Problem: sequential input

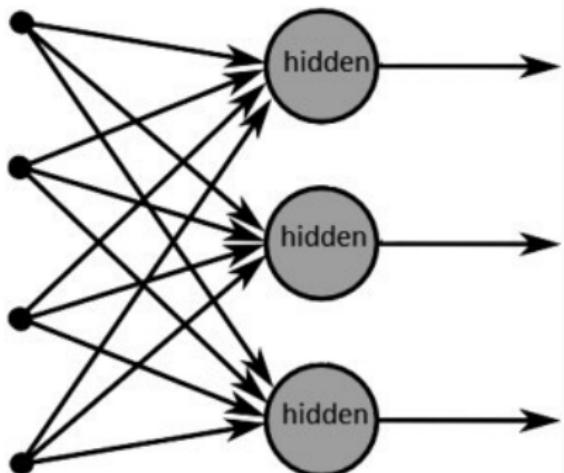


- What if we want to do machine learning on sequential data (i.e. music, text, time series, etc.)?
- How do we feed in a sequence of inputs into a neural network?
- How does a neural network keep track of all of the information it receives?

Solution: recurrence!



(a) Recurrent neural network



(b) Forward neural network

The deep learning revolution



Necip Fazil Ayan

1 hr ·

Old:

Onlarin, İzmir'in neden hayır dediğini anlamalarını beklemiyoruz.

Their, Izmir's why you said no we don't expect them to understand.



• Rate this translation



Necip Fazil Ayan

1 hr ·

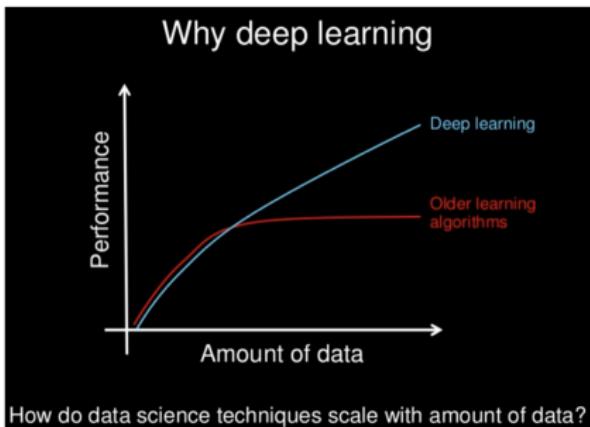
New:

Onlarin, İzmir'in neden hayır dediğini anlamalarını beklemiyoruz.

We don't expect them to understand why Izmir said no.

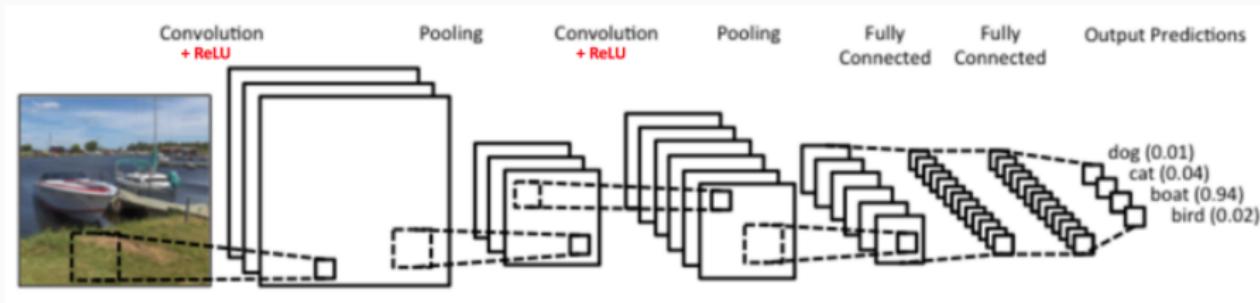


• Rate this translation

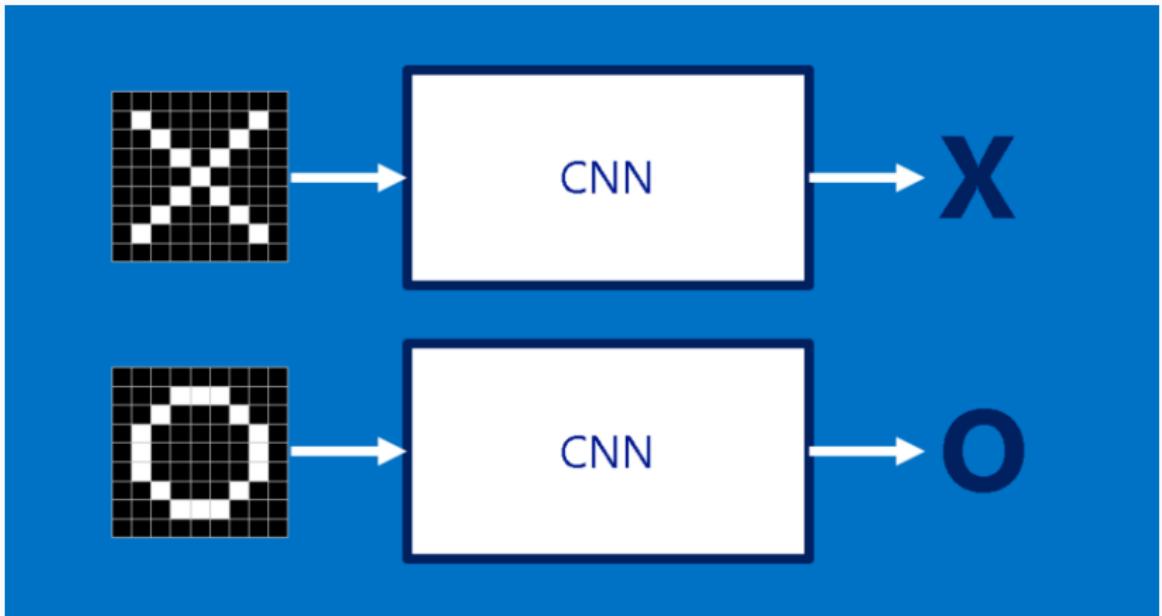


Convolutional Neural Networks

Our first CNN



Let's look at a toy example!



- How would a CNN classify X's vs O's?

Would you classify this as an X?

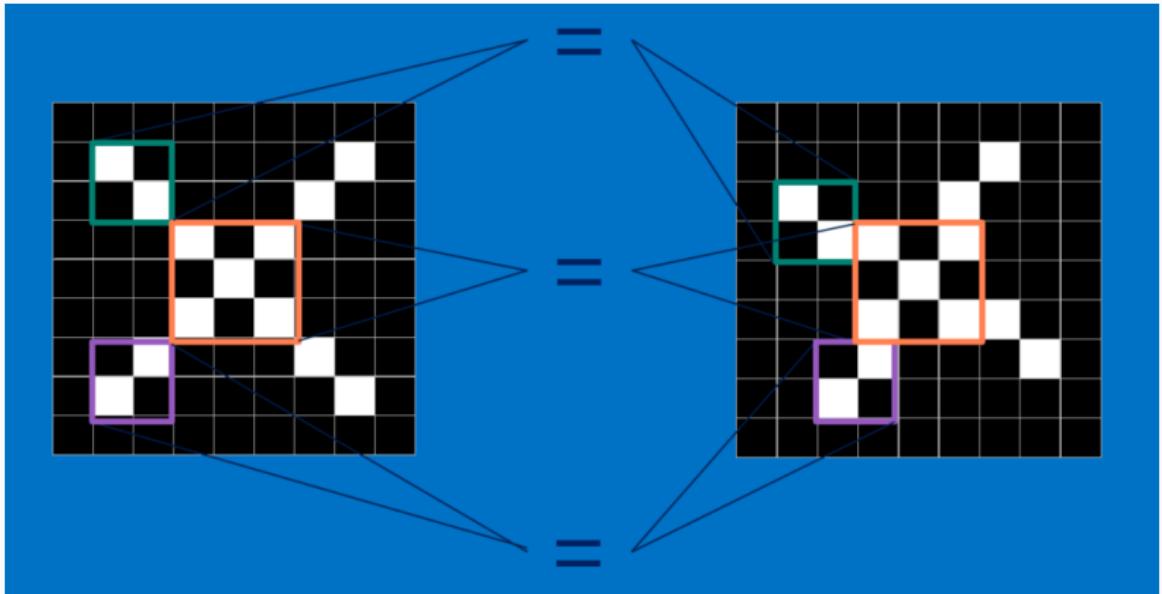
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	-1	-1	-1	-1	-1	-1	-1	



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

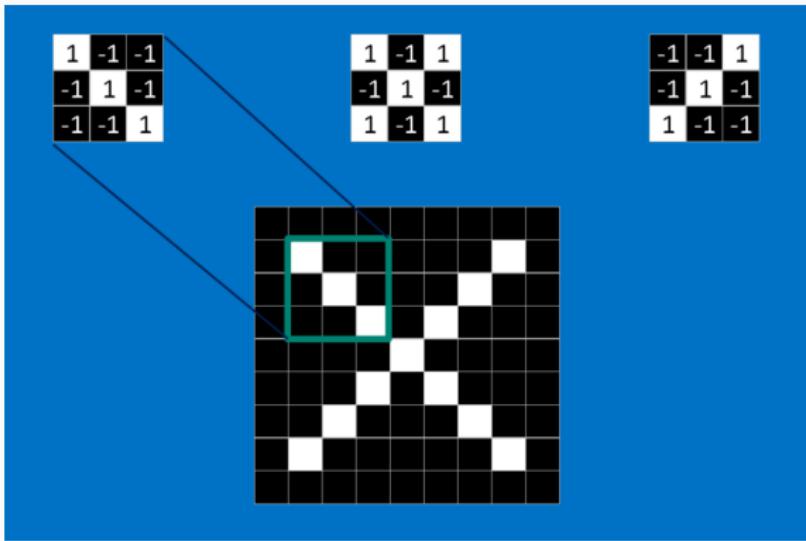
- A computer can't easily distinguish this example as an X.. so what can we do?

What is a feature?



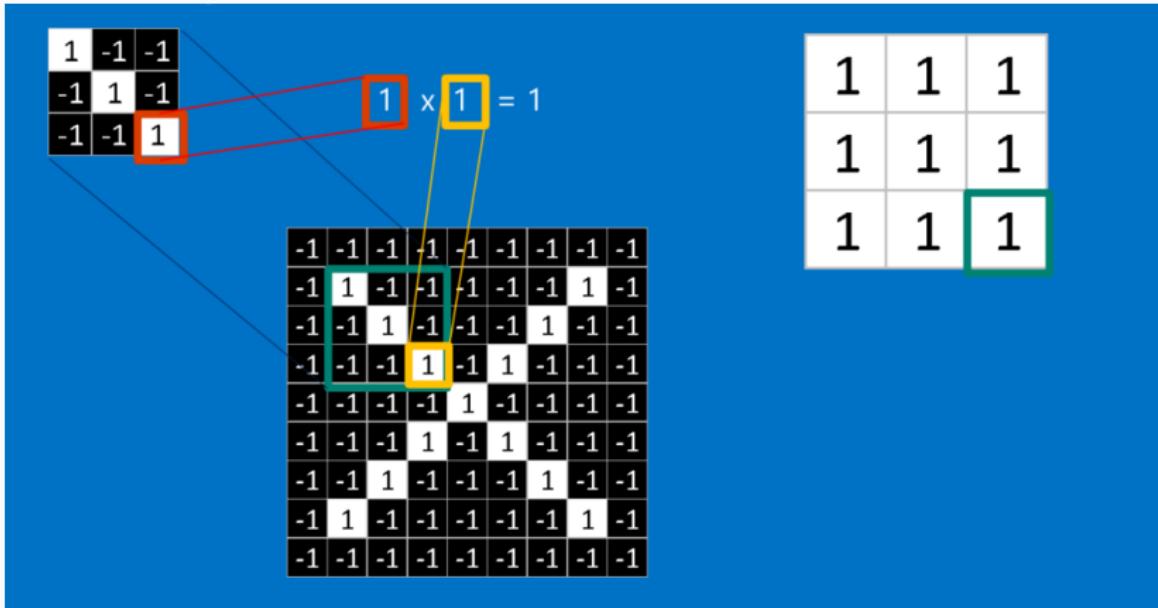
- We can break down the image into smaller images and look for patterns!
- Can we see here the similar patterns?
- These diagonals and crosses can be thought of as features!

Filters



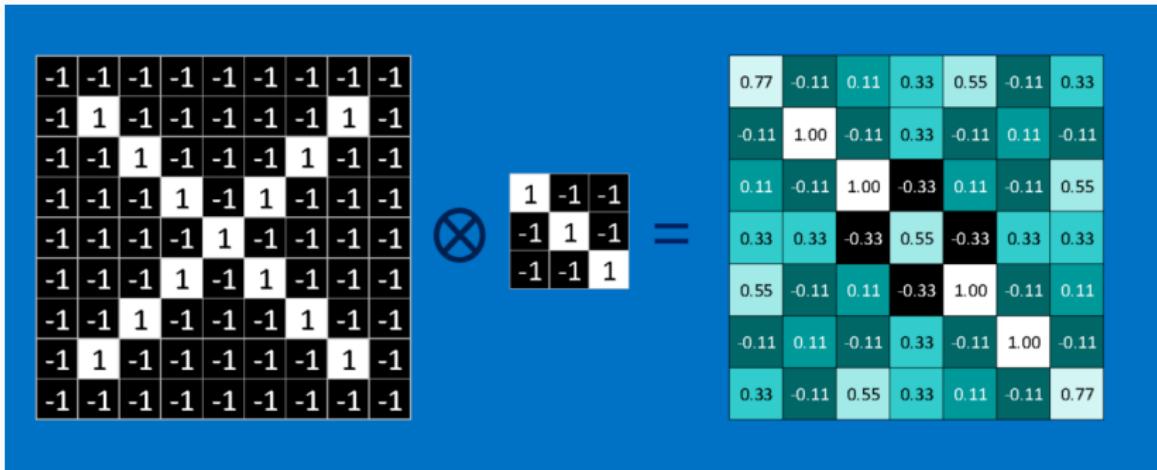
- Let's break this up into smaller problems... what about smaller patterns!

CNN Layers: Convolutions



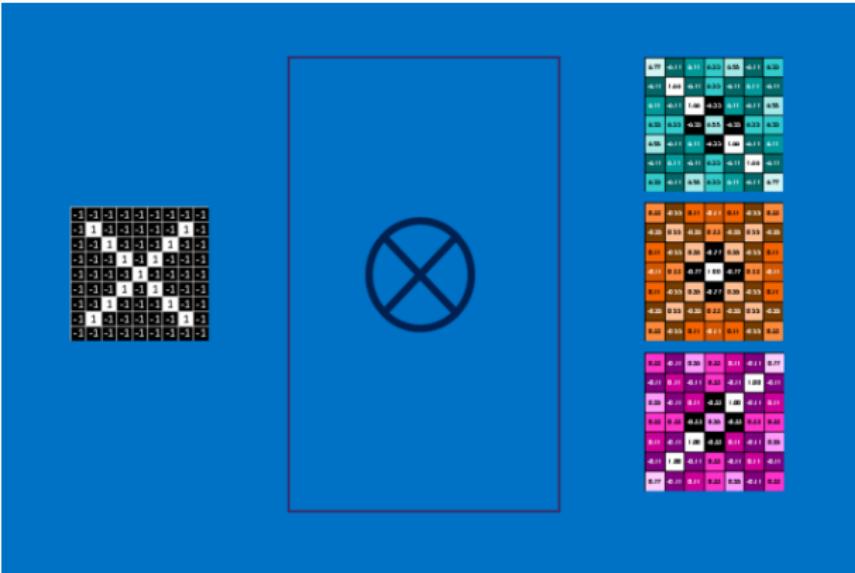
- Now let's compare our smaller pattern with the image.
- This is the basic idea behind convolutions!

CNN Layers: Convolutions



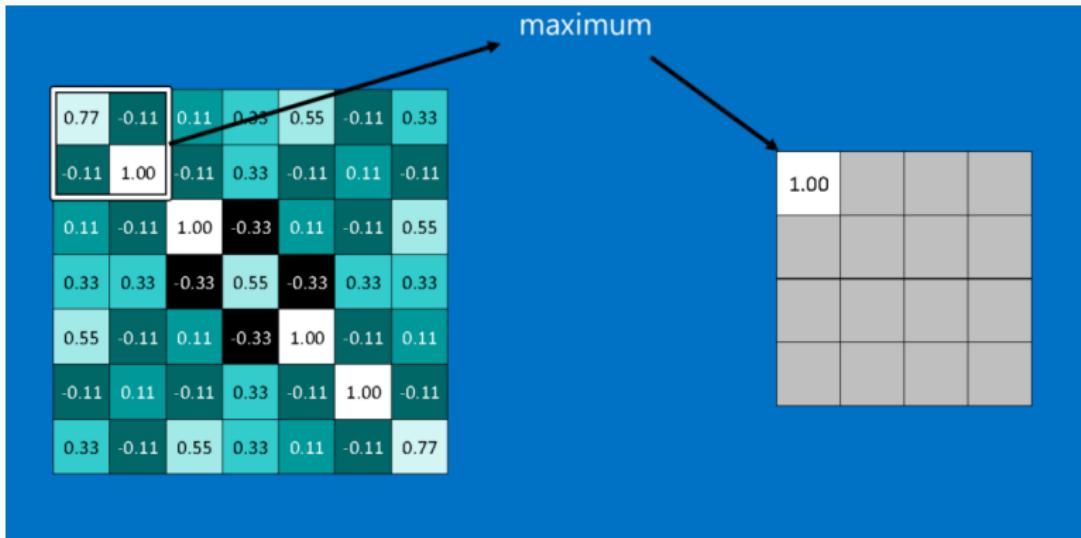
- Now let's apply this process to every part of the board.
- We can look at the new board and call it a filtered version of this specific feature
- Shows us where on the image the feature matches the most

CNN Layers: Convolutions



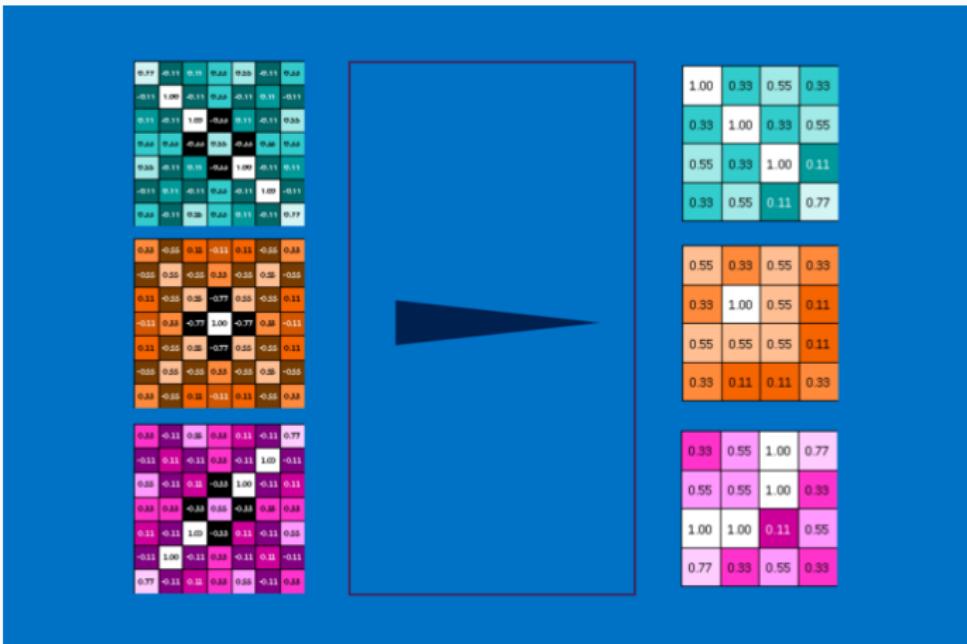
- Now lets apply this process to every part of the image for different features.

CNN Layers: Pooling



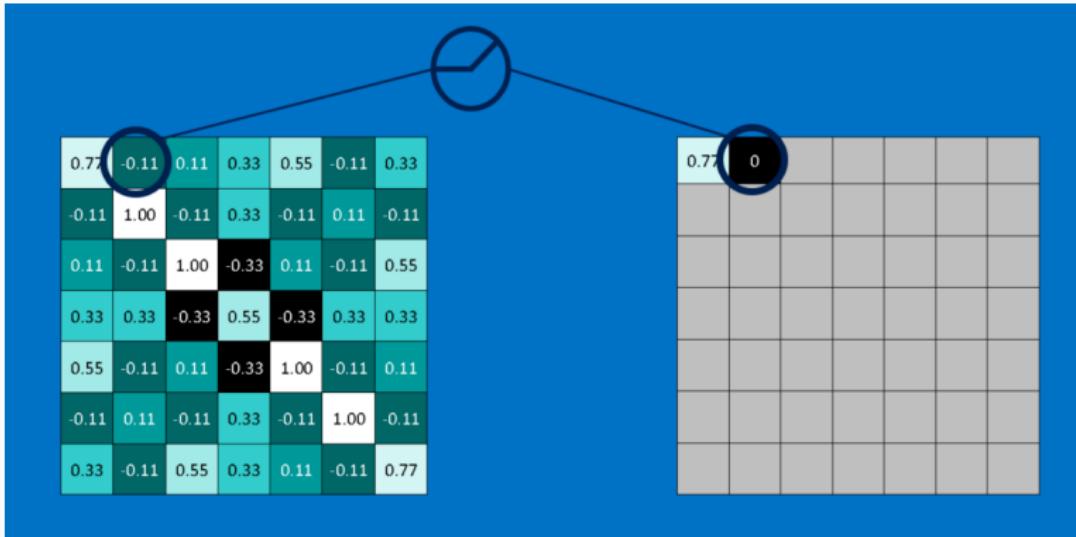
- If CNN's are really deep, doing that math for every possible pixel gets to be quite computationally expensive
- That's why we have pooling layers, that down sample our image!

CNN Architecture



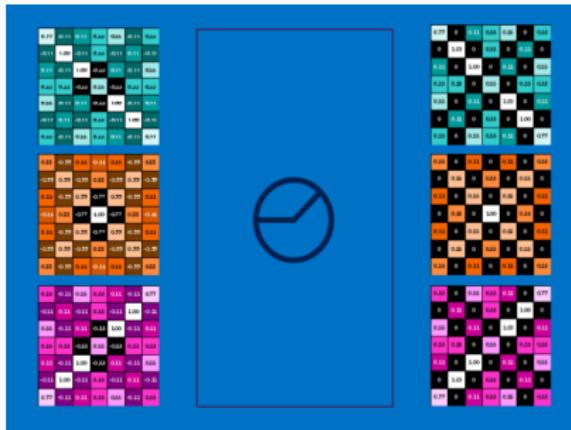
- By applying these pooling layers, we try to capture the patterns and get rid of extra unneeded information.

CNN Layers: Activation



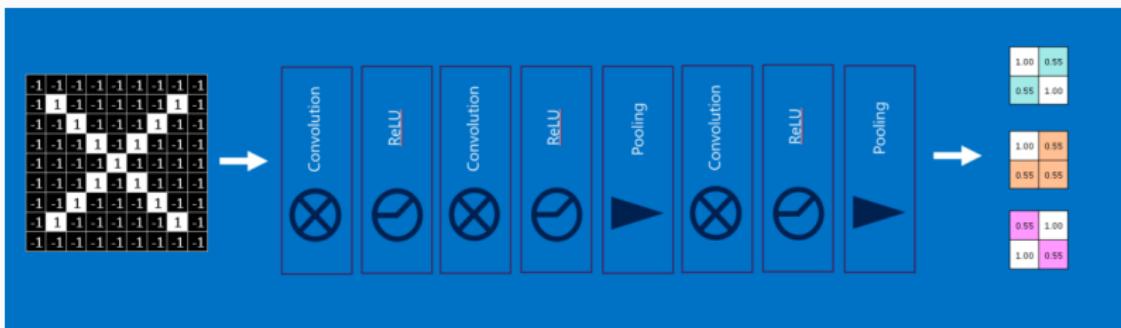
- Remember those activation functions you learned from the last lecture? We use them here too!
- Applying RELU we simply get rid of all the negatives!

CNN Pooling Layers

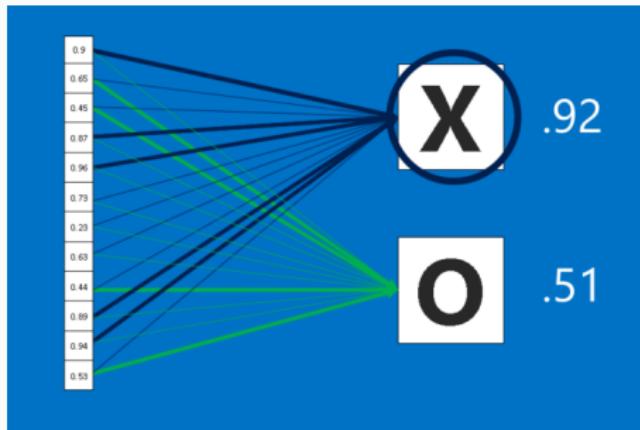


- The idea behind this, is to help keep any one number stuck at 0, or block up towards infinity
 - This may not seem like a big issue in our toy example, but in much larger more complicated CNN's having stable values is important!

Let's put it all together!



CNN: Fully Connected Layer



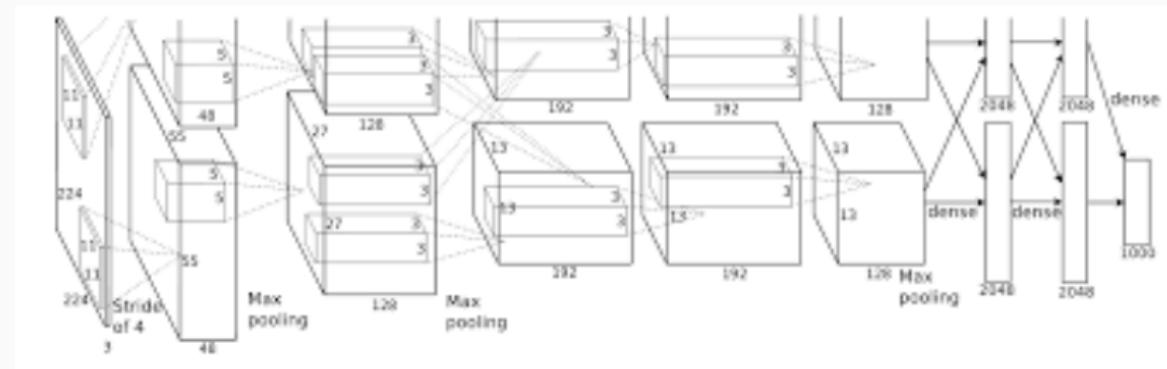
- At the end we use a fully connected layer to hold a vote if the image is either an X or O
- Each value has a weight though, so not all values contribute equally
- This goes back to the ideas you learned last week!

Applications of CNNs

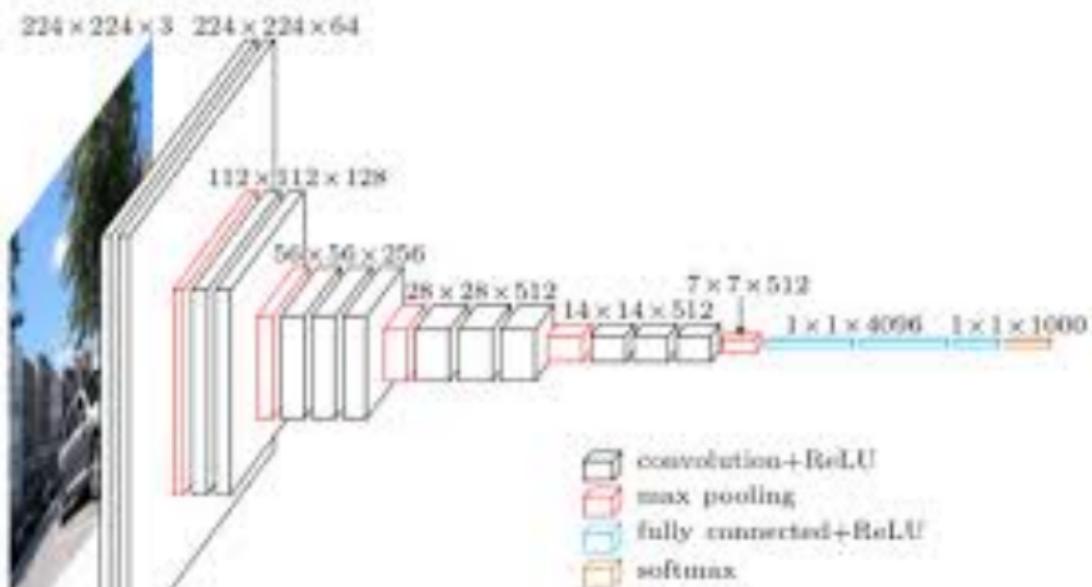
- Computer vision
 - Face recognition
 - Scene labeling
 - Image segmentation
 - Image classification
 - Action recognition
 - Human pose estimation
 - Document analysis
- Natural language processing
 - Speech recognition
 - Text classification
- Games (Checkers, Go)
- Drug discovery
- Medical image analysis
- Aircraft ejection seat testing

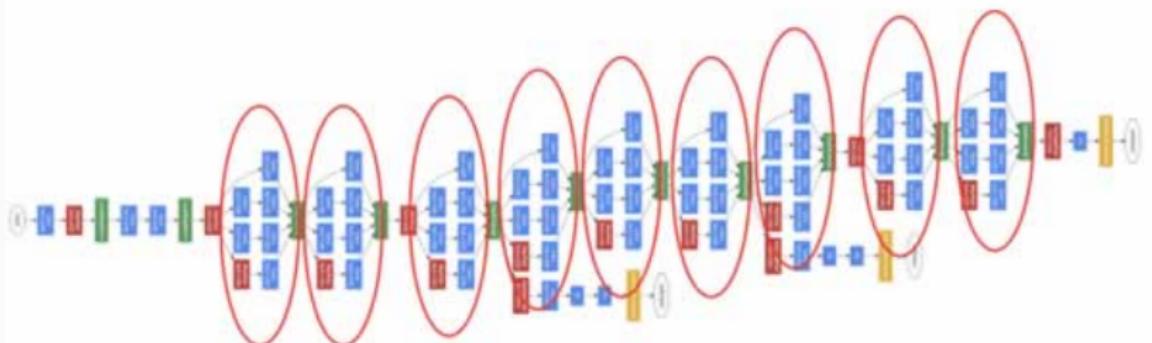
- "The 9 Deep Learning Papers You Need to Know About" (<https://adeshpande3.github.io/adeshpande3.github.io/The-9-Deep-Learning-Papers-You-Need-To-Know-About.html>)
- Many of these models are used as benchmarks in deep learning
- There exist many pre-trained versions of the above models that can be leveraged for a new task (this is called **transfer learning**)
- Let's look at a few...

AlexNet (2012)



VGG-16 (2014)



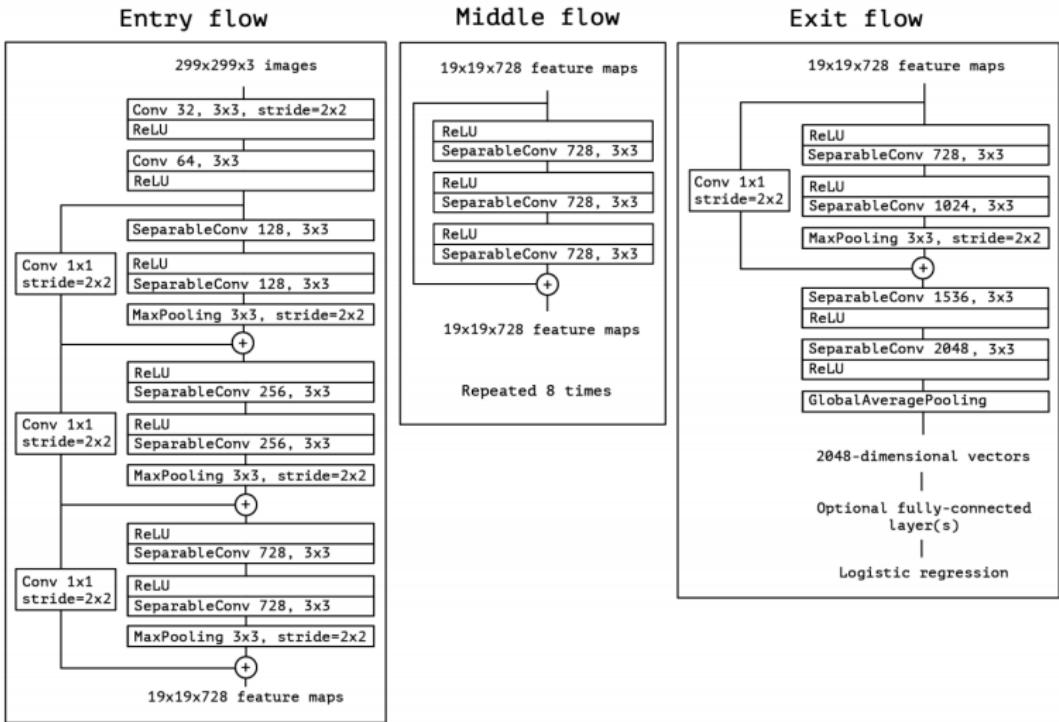


9 Inception modules

Network in a network in a network...

Convolution
Pooling
Softmax
Other

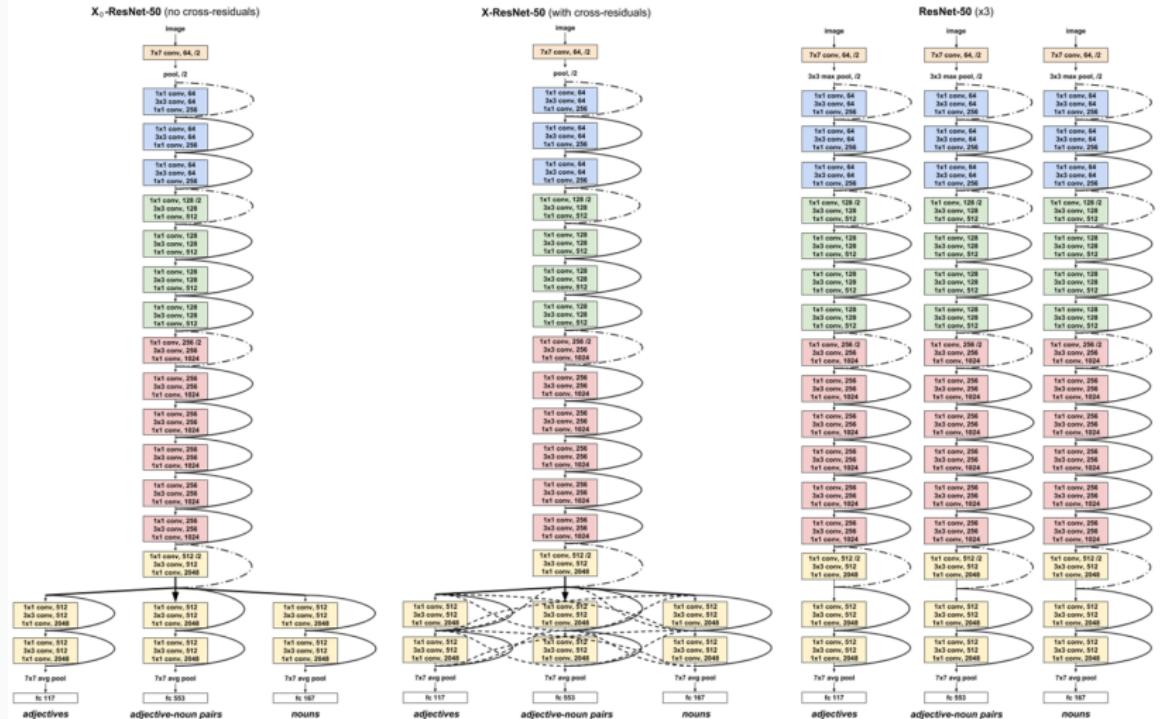
Xception (2016)



ResNet-50 (2015)

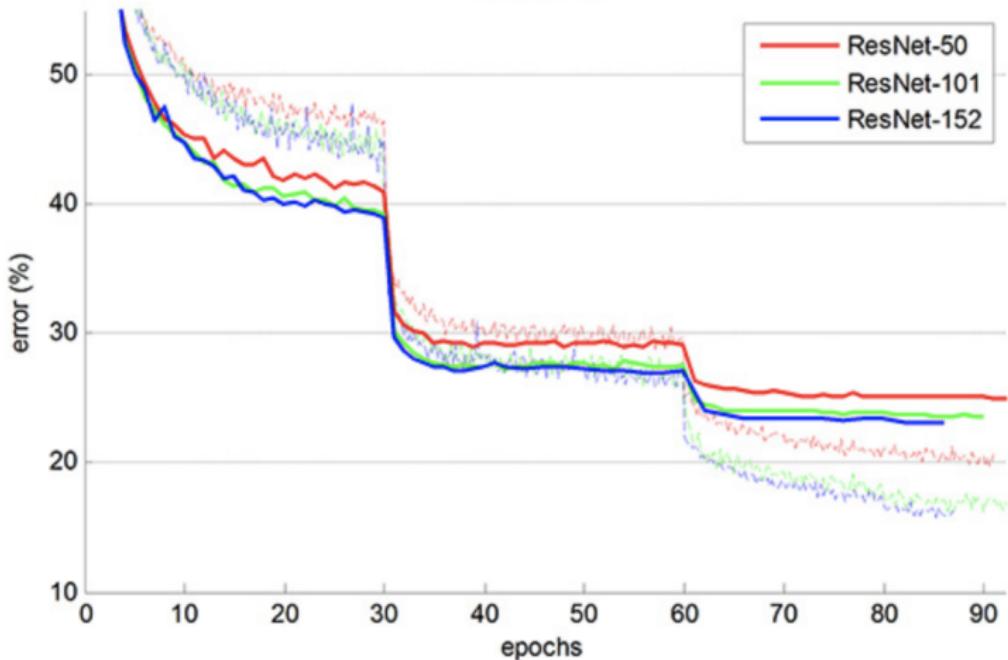


ML@B



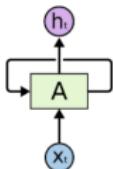
ResNet-152 (2015)

ImageNet-1k



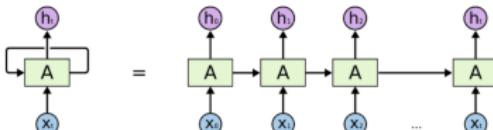
Recurrent Neural Networks

- RNNs are networks with loops, allowing information over time to persist



Recurrent Neural Networks have loops.

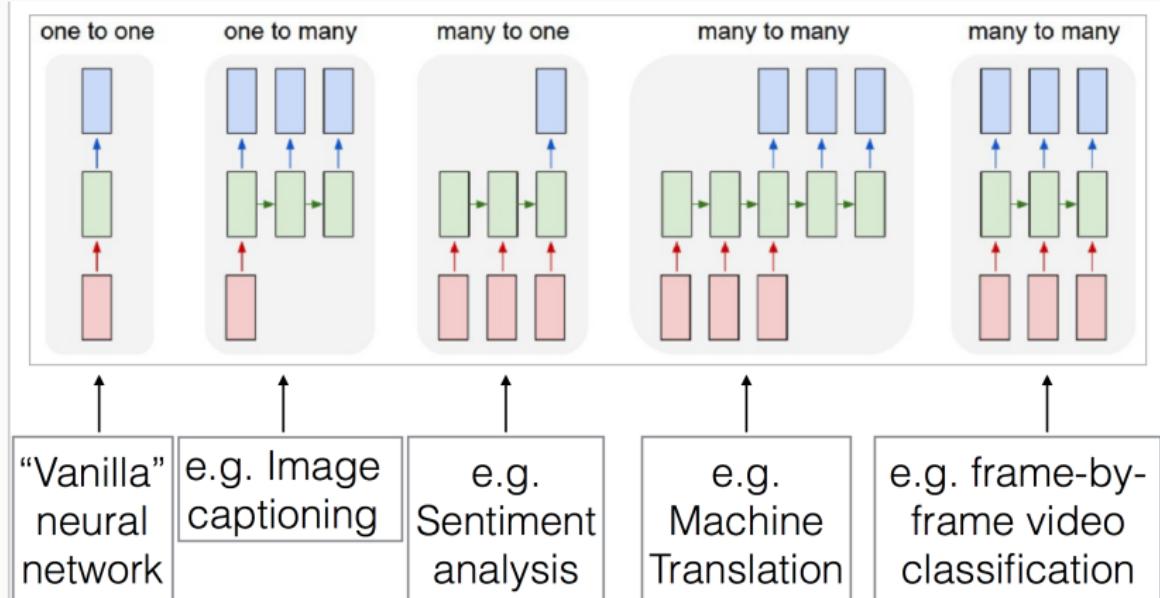
- Consider unrolling the loop:



An unrolled recurrent neural network.

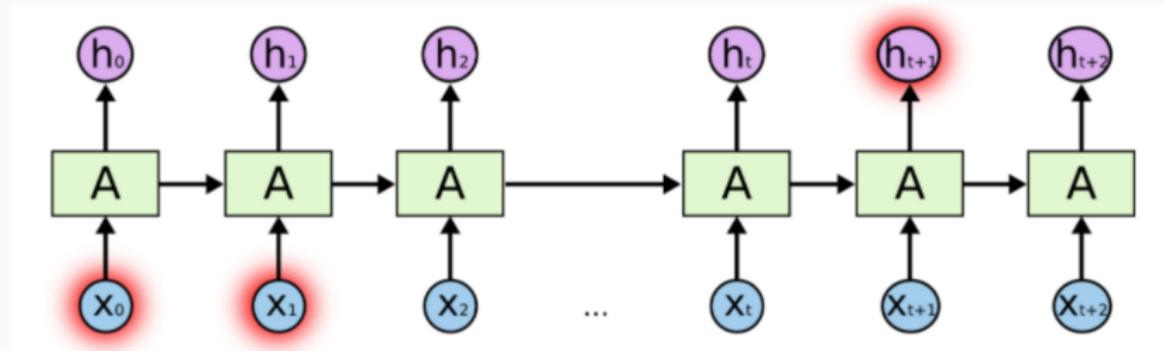
²<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Examples of recurrent neural networks



³<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Problem: long-term dependencies



"I grew up in *France*... I speak fluent *French*"

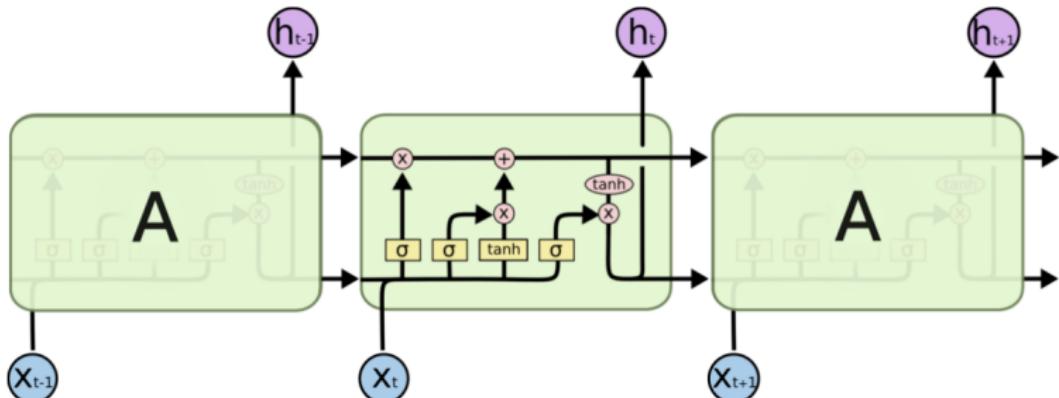
⁴<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Solution: LSTMs



- "Long Short Term Memory Network"
- Introduced by Hochreiter & Schmidhuber in 1997
- Capable of learning long-term dependencies
- Instead of a single neural network layer, there are 4 which interact in a special way

Inside of an LSTM cell



Neural Network Layer

Pointwise Operation

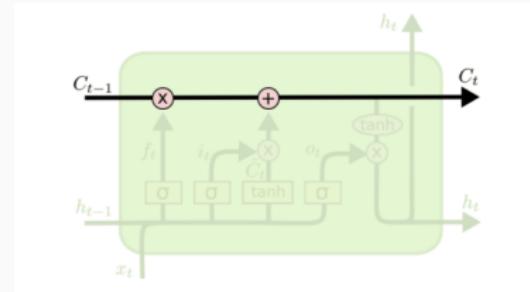
Vector Transfer

Concatenate

Copy

⁵<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

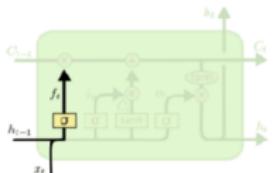
Core idea behind LSTMs



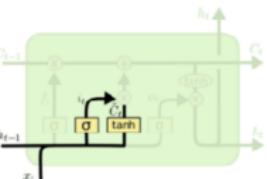
- The cell state is like a conveyor belt on which information flows
- The LSTM can add or remove information from the cell state, regulated by structures called gates
- Gates are made of a sigmoid neural net layer and a pointwise multiplication operation.

⁶<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Steps in an LSTM

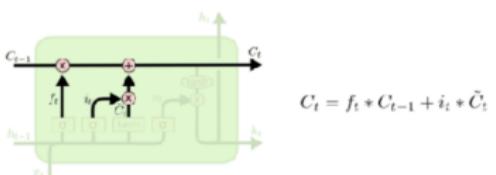


$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$



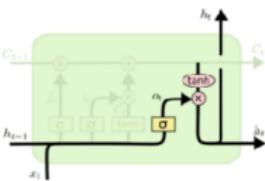
$$\begin{aligned} i_t &= \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \\ \tilde{C}_t &= \tanh(W_C \cdot [h_{t-1}, x_t] + b_C) \end{aligned}$$

1. Input gate



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

3. Update cell state



$$\begin{aligned} o_t &= \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) \\ h_t &= o_t * \tanh(C_t) \end{aligned}$$

4. Output new hidden state and new cell state

⁷<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Applications of LSTMs



- Time series prediction (finance, economics, etc.)
- Speech recognition
- Sentiment analysis
- Machine translation
- Music composition
- Robot control
- Handwriting recognition
- Human action recognition
- Protein homology detection
- Predicting subcellular localization of proteins
- Text generation
- Many, many more...

- Training networks using GPUs
- CNN-RNN combinations
- Network visualization and interpretation
- Transfer learning (see project 3!)
- Other types of neural networks: auto-encoders, Neural Turing machines, Generative adversarial networks, ...
- The future of all of this

Questions

Questions?

Demos

CNN demo: FashionMNIST

RNN demo: sentiment classification